

Jason wants to marry Kriti. But Kriti's father says Jason first needs to be competent enough to marry his daughter. So he tells Jason that he needs to buy a house, a car and a computer (Kriti is a binge-queen) to be worthy offer. Love knows no bounds and neither does Jason. So he decides to take up this challenge and works day and night for the next 1 year. After toiling away for such a long time, he now needs to calculate the total amount of money he requires to fulfill Kriti's (but really her father's) wishes to find out if he has enough saved up right now or not. Data has been given below and you need to use that to help him figure out how much he would need to spend to afford all of the above mentioned things. Your goal is to give Jason a total estimate of his expenses.

In []:

```
pip install category_encoders
```

```
Requirement already satisfied: category_encoders in /usr/local/lib/python3.7/dist-packages (2.2.2)
Requirement already satisfied: numpy>=1.14.0 in /usr/local/lib/python3.7/dist-packages (from category_encoders) (1.19.5)
Requirement already satisfied: statsmodels>=0.9.0 in /usr/local/lib/python3.7/dist-packages (from category_encoders) (0.10.2)
Requirement already satisfied: scipy>=1.0.0 in /usr/local/lib/python3.7/dist-packages (from category_encoders) (1.4.1)
Requirement already satisfied: pandas>=0.21.1 in /usr/local/lib/python3.7/dist-packages (from category_encoders) (1.1.5)
Requirement already satisfied: scikit-learn>=0.20.0 in /usr/local/lib/python3.7/dist-packages (from category_encoders) (0.22.2.post1)
Requirement already satisfied: patsy>=0.5.1 in /usr/local/lib/python3.7/dist-packages (from category_encoders) (0.5.1)
Requirement already satisfied: pytz>=2017.2 in /usr/local/lib/python3.7/dist-packages (from pandas>=0.21.1->category_encoders) (2018.9)
Requirement already satisfied: python-dateutil>=2.7.3 in /usr/local/lib/python3.7/dist-packages (from pandas>=0.21.1->category_encoders) (2.8.1)
Requirement already satisfied: joblib>=0.11 in /usr/local/lib/python3.7/dist-packages (from scikit-learn>=0.20.0->category_encoders) (1.0.1)
Requirement already satisfied: six in /usr/local/lib/python3.7/dist-packages (from patsy>=0.5.1->category_encoders) (1.15.0)
```

Necessary Modules :

In []:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn import preprocessing, svm, model_selection
from sklearn.linear_model import LinearRegression, Lasso, Ridge
import seaborn as sns
import category_encoders as ce
```

```
/usr/local/lib/python3.7/dist-packages/statsmodels/tools/_testing.py:19: FutureWarning: pandas.util.testing is deprecated. Use the functions in the public API at pandas.testing instead.
```

```
import pandas.util.testing as tm
```

Reading data :

In []:

```
data=pd.read_csv('laptops.csv',encoding='ISO-8859-1',index_col=0)
```

In []:

```
data
```

Out []:

Company		Product	TypeName	Inches	ScreenResolution	Cpu	Ram	Memory	Gpu	OpSys
1	Apple	MacBook Pro	Ultrabook	13.3	IPS Panel Retina Display 2560x1600	Intel Core i5 2.3GHz	8GB	128GB SSD	Intel Iris Plus Graphics 640	macOS
2	Apple	Macbook Air	Ultrabook	13.3	1440x900	Intel Core i5 1.8GHz	8GB	128GB Flash Storage	Intel HD Graphics 6000	macOS
3	HP	250 G6	Notebook	15.6	Full HD 1920x1080	Intel Core i5 7200U 2.5GHz	8GB	256GB SSD	Intel HD Graphics 620	No OS
4	Apple	MacBook Pro	Ultrabook	15.4	IPS Panel Retina Display 2880x1800	Intel Core i7 2.7GHz	16GB	512GB SSD	AMD Radeon Pro 455	macOS
5	Apple	MacBook Pro	Ultrabook	13.3	IPS Panel Retina Display 2560x1600	Intel Core i5 3.1GHz	8GB	256GB SSD	Intel Iris Plus Graphics 650	macOS
...
1316	Lenovo	Yoga 500-14ISK	2 in 1 Convertible	14.0	IPS Panel Full HD / Touchscreen 1920x1080	Intel Core i7 6500U 2.5GHz	4GB	128GB SSD	Intel HD Graphics 520	Windows 10
1317	Lenovo	Yoga 900-13ISK	2 in 1 Convertible	13.3	IPS Panel Quad HD+ / Touchscreen 3200x1800	Intel Core i7 6500U 2.5GHz	16GB	512GB SSD	Intel HD Graphics 520	Windows 10
1318	Lenovo	IdeaPad 100S-14IBR	Notebook	14.0	1366x768	Intel Celeron Dual Core N3050 1.6GHz	2GB	64GB Flash Storage	Intel HD Graphics	Windows 10
1319	HP	15-AC110nv (i7-6500U/6GB/1TB/Radeon)	Notebook	15.6	1366x768	Intel Core i7 6500U 2.5GHz	6GB	1TB HDD	AMD Radeon R5 M330	Windows 10
1320	Asus	X553SA-XX031T (N3050/4GB/500GB/W10)	Notebook	15.6	1366x768	Intel Celeron Dual Core N3050 1.6GHz	4GB	500GB HDD	Intel HD Graphics	Windows 10

1303 rows × 12 columns



In []:

```
data.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 1303 entries, 1 to 1320
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Company               1303 non-null   object
1   Product               1303 non-null   object
2   TypeName              1303 non-null   object
3   Inches                1303 non-null   float64
4   ScreenResolution      1303 non-null   object
5   Cpu                   1303 non-null   object
6   Ram                   1303 non-null   object
7   Memory                1303 non-null   object
8   Gpu                   1303 non-null   object
9   OpSys                 1303 non-null   object
```

```
10  Weight          1303 non-null  object
11  Price_euros      1303 non-null  float64
dtypes: float64(2), object(10)
memory usage: 132.3+ KB
```

Check for any missing values that are present in the dataset :

```
In [ ]:
```

```
data.isnull().sum()
```

```
Out[ ]:
```

```
Company          0
Product          0
TypeName         0
Inches           0
ScreenResolution 0
Cpu              0
Ram              0
Memory           0
Gpu              0
OpSys            0
Weight           0
Price_euros      0
dtype: int64
```

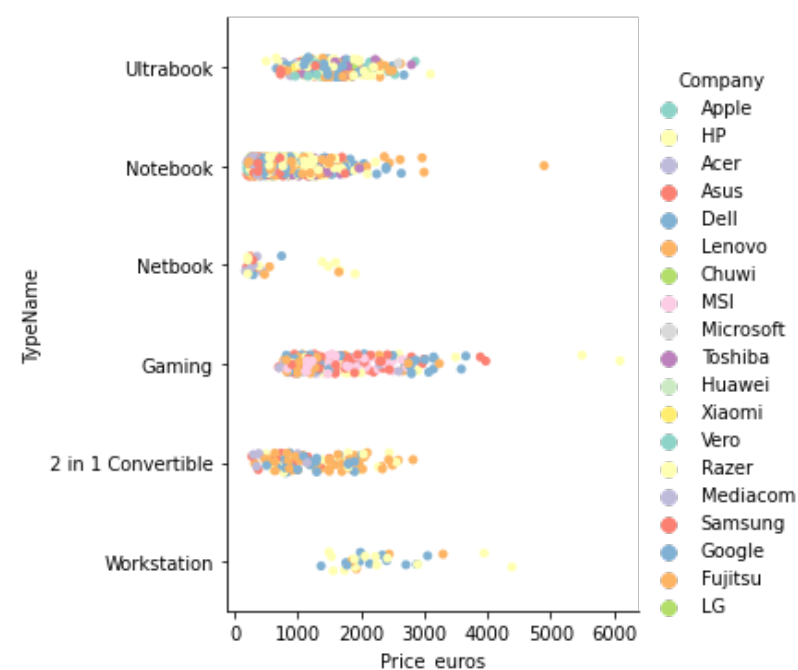
Plotting :

```
In [ ]:
```

```
sns.catplot(x = "Price_euros", y = 'TypeName', data = data, hue = "Company", palette="Set3")
```

```
Out[ ]:
```

```
<seaborn.axisgrid.FacetGrid at 0x7fa7d0ecdb90>
```





```
In [ ]:
```

```
sns.catplot(x = "Price_euros", y = 'Ram', data = data, hue = "Company", palette="Set3")
```

```
Out[ ]:
```

```
<seaborn.axisgrid.FacetGrid at 0x7fa7bc23c790>
```



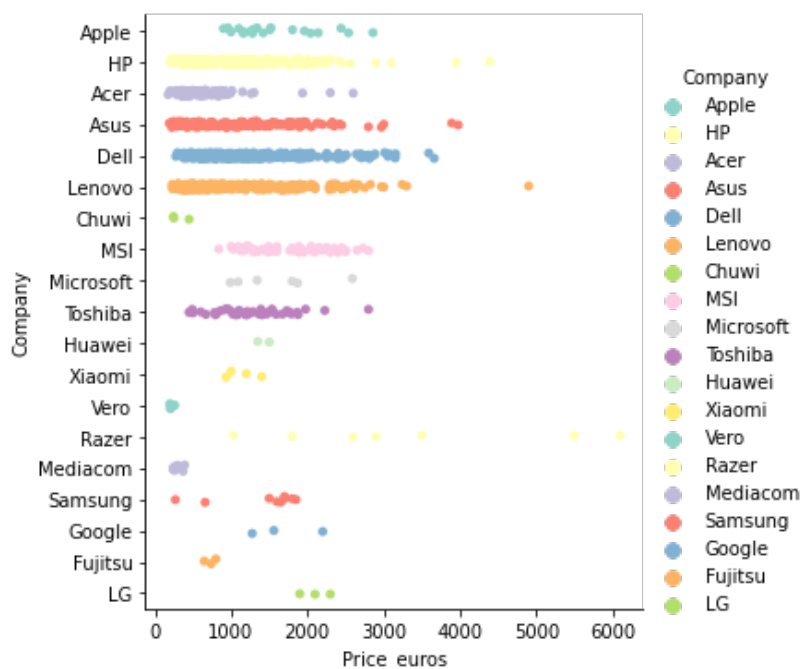


In []:

```
sns.catplot(x = "Price_euros", y = 'Company', data = data, hue = "Company", palette="Set3")
```

Out[]:

<seaborn.axisgrid.FacetGrid at 0x7fa7a880d890>





In []:

```
plt.figure(figsize=(15,10))
sns.set_theme(style="whitegrid")
sns.boxplot(x='TypeName', y='Price_euros', data = data, palette='gist_rainbow', linewidth=5)
```

Out[]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f19eced8510>





Dealing with outliers :

In []:


```
Q1 = data['Price_euros'].quantile(0.25)
Q3 = data['Price_euros'].quantile(0.75)
IQR = Q3 - Q1
filter = (data['Price_euros'] >= Q1 - 1.5 * IQR) & (data['Price_euros'] <= Q3 + 1.5 * IQR)
data = data.loc[filter]
```

In []:

```
plt.figure(figsize=(15,10))
sns.set_theme(style="whitegrid")
sns.boxplot(x='TypeName',y='Price_euros',data = data,palette='gist_rainbow',linewidth=5)
```

Out[]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f19eb0e17d0>



Process of Data Cleaning:

In []:

```
def sufrem(clmn):
    clmn = clmn.str.replace(r'TB', '')
    clmn = clmn.str.replace(r'GB', '').fillna(0)
    return clmn
def TBhandle(clmn):
    if clmn < 6 :
        return clmn*1024
    else :
        return clmn
```

In []:

```
def clean(df):
    df['Ram'] = df['Ram'].replace('[GB]', '', regex=True)
    df['Ram'] = pd.to_numeric(df['Ram'])
    df['Weight'] = df['Weight'].replace('[kg]', '', regex=True)
    df['Weight'] = pd.to_numeric(df['Weight'])
    df['cpu_speed(GHz)'] = pd.to_numeric(df['Cpu'].str.split(' ').str[-1].replace('[GHz]', '', regex=True))
    df['cpu_manuf'] = df['Cpu'].str.split(' ').str[0]
    df['Horizontal'] = pd.to_numeric(df['ScreenResolution'].str.split(' ').str[-1].str.split('x').str[0])
    df['Vertical'] = pd.to_numeric(df['ScreenResolution'].str.split(' ').str[-1].str.split('x').str[1])
    t_arr = []
    for i in df['ScreenResolution']:
        s = i.split(' ')
        if "Touchscreen" in s:
            touch = True
        else:
            touch = False
        t_arr.append(touch)
    df['Touch'] = pd.to_numeric(t_arr)
    SSD = []
    HDD = []
    Hybrid = []
    Flash = []
    for i in df['Memory']:
        dual_s = i.split(' ')
        if "SSD" in dual_s:
            SSD_val = dual_s[dual_s.index('SSD')-1]
            SSD.append(SSD_val)
        else:
            SSD.append('0')
        if "HDD" in dual_s:
            HDD_val = dual_s[dual_s.index("HDD")-1]
            HDD.append(HDD_val)
        else:
            HDD.append('0')
        if "Flash" in dual_s:
            Flash_val = dual_s[dual_s.index("Flash")-1]
            Flash.append(Flash_val)
        else:
            Flash.append('0')
        if "Hybrid" in dual_s:
            Hybrid_val = dual_s[dual_s.index("Hybrid")-1]
            Hybrid.append(Hybrid_val)
        else:
            Hybrid.append('0')
    df['Flash'] = Flash
    df['HDD'] = HDD
    df['SSD'] = SSD
    df['Hybrid'] = Hybrid
    df[['HDD', 'SSD', 'Flash', 'Hybrid']] = df[['HDD', 'SSD', 'Flash', 'Hybrid']].apply(sufrem)

    df['HDD'] = pd.to_numeric(df['HDD']).apply(TBhandle)
    df['SSD'] = pd.to_numeric(df['SSD']).apply(TBhandle)
```

```

df['Flash'] = pd.to_numeric(df['Flash']).apply(TBhandle)
df['Hybrid'] = pd.to_numeric(df['Hybrid']).apply(TBhandle)
cpu_dummies = pd.get_dummies(df['cpu_manuf'], prefix = 'cpu')
opsys_dummies = pd.get_dummies(df['OpSys'], prefix = 'op')
type_dummies = pd.get_dummies(df['TypeName'], prefix = 'Type')
df = pd.concat([cpu_dummies, type_dummies, opsys_dummies, df], axis=1)
df.drop('cpu_manuf',axis = 1, inplace = True)
df.drop('OpSys',axis = 1, inplace = True)
df.drop('TypeName',axis = 1, inplace = True)
df.drop(columns=['Memory', 'Company', 'Product', 'ScreenResolution', 'Cpu'], inplace = True)

return df

df = clean(data)

```

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

This is separate from the ipykernel package so we can avoid doing imports until
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:4: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

after removing the cwd from sys.path.
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:5: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

"""
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:6: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:7: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
import sys

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:8: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:9: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

if __name__ == '__main__':
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:18: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:45: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:46: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:47: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:48: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

/usr/local/lib/python3.7/dist-packages/pandas/core/frame.py:3069: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

self[k1] = value[k2]

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:50: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:51: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:52: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:53: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

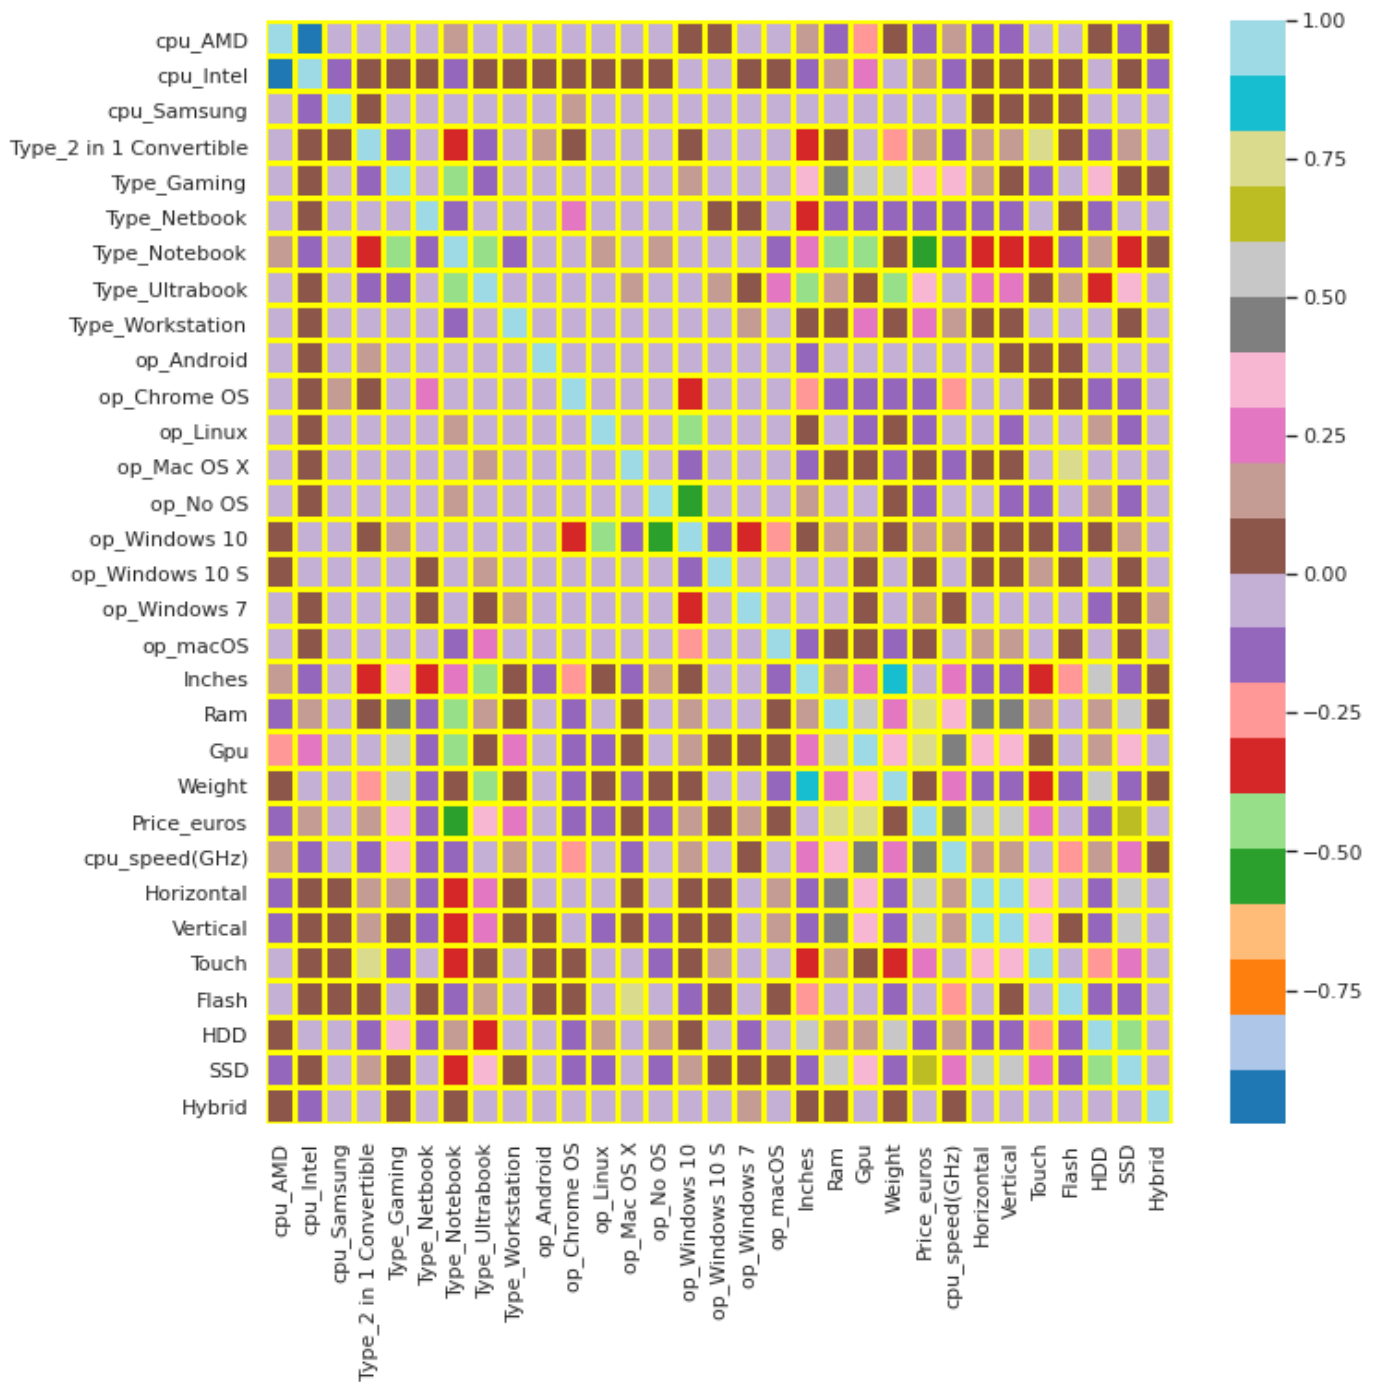
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

In []:

```
corrmat = df.corr()
fig, ax = plt.subplots()
fig.set_size_inches(11,11)
sns.heatmap(corrmat,cmap = "tab20",linewidths = 2,linecolor = "yellow")
```

Out[]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f19e4b37c90>



Target Encoding :

In []:

```
encoder=ce.TargetEncoder(cols='Gpu')
gpu_encoded = encoder.fit_transform(df['Gpu'],df['Price_euros'])
df['Gpu'] = gpu_encoded['Gpu']
```

/usr/local/lib/python3.7/dist-packages/category_encoders/utils.py:21: FutureWarning: is_categorical is deprecated and will be removed in a future version. Use is_categorical_dtype instead

```
elif pd.api.types.is_categorical(cols):
```

In []:

```
temp1 = df['Gpu']
temp2 = data['Gpu']
gpudf = pd.DataFrame()
gpudf['Gpu'] = temp2
gpudf['val'] = temp1
```

In []:

```
df.head()
```

Out []:

	cpu_AMD	cpu_Intel	cpu_Samsung	Type_2 in 1 Convertible	Type_Gaming	Type_Netbook	Type_Notebook	Type_Ultrabook	Type_Work
1	0	1	0	0	0	0	0	1	
2	0	1	0	0	0	0	0	1	
3	0	1	0	0	0	0	1	0	
4	0	1	0	0	0	0	0	1	
5	0	1	0	0	0	0	0	1	

In []:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1274 entries, 1 to 1320
Data columns (total 31 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   cpu_AMD                             1274 non-null   uint8
1   cpu_Intel                           1274 non-null   uint8
2   cpu_Samsung                         1274 non-null   uint8
3   Type_2 in 1 Convertible             1274 non-null   uint8
4   Type_Gaming                         1274 non-null   uint8
5   Type_Netbook                       1274 non-null   uint8
6   Type_Notebook                      1274 non-null   uint8
7   Type_Ultrabook                     1274 non-null   uint8
8   Type_Workstation                    1274 non-null   uint8
9   op_Android                         1274 non-null   uint8
10  op_Chrome OS                       1274 non-null   uint8
11  op_Linux                           1274 non-null   uint8
12  op_Mac OS X                       1274 non-null   uint8
13  op_No OS                          1274 non-null   uint8
14  op_Windows 10                     1274 non-null   uint8
15  op_Windows 10 S                   1274 non-null   uint8
16  op_Windows 7                      1274 non-null   uint8
17  op_macOS                          1274 non-null   uint8
18  Inches                             1274 non-null   float64
19  Ram                                1274 non-null   int64
20  Gpu                                1274 non-null   float64
21  Weight                             1274 non-null   float64
22  Price_euros                        1274 non-null   float64
23  cpu_speed(GHz)                    1274 non-null   float64
24  Horizontal                         1274 non-null   int64
25  Vertical                           1274 non-null   int64
26  Touch                              1274 non-null   bool
27  Flash                              1274 non-null   int64
28  HDD                                1274 non-null   float64
29  SSD                                1274 non-null   int64
30  Hybrid                             1274 non-null   float64
dtypes: bool(1), float64(7), int64(5), uint8(18)
memory usage: 193.0 KB
```

Building a Model :

In []:

```
y = (df['Price_euros'])
X = np.array(df.drop(['Price_euros'],1))
```

In []:

```
X_train , X_test , y_train , y_test = model_selection.train_test_split(X , y , test_size
= 0.2, random_state = 22)
```

Cleaning Data To Be Predicted :

In []:

```
pred_df = pd.DataFrame({'Company':['HP'],'Product':['250 G6'],'TypeName' : ['Notebook'],  
'Inches' : [14],'ScreenResolution' : ['IPS Panel Full HD 1920x1080'], 'Cpu' : ['Intel Co  
re i7 7500U 2.7GHz'],'Ram' : ['8GB'],'Memory' : ['1TB SSD'],'Gpu' : ['Intel HD Graphics  
620'], 'OpSys' : ['Linux'], 'Weight' : ['1.5kg']})  
pred_df = clean(pred_df)
```

In []:

```
l = ['cpu_AMD','cpu_Intel','cpu_Samsung','Type_2 in 1 Convertible','Type_Gaming','Type_Ne  
tbook','Type_Notebook','Type_Ultrabook','Type_Workstation','op_Android','op_Chrome', 'op_  
Linux',  
      'op_Mac OS X', 'op_No OS', 'op_Windows 10', 'op_Windows 10 S',  
      'op_Windows 7', 'op_macOS']  
for col in l:  
    pred_df[col] = 0
```

In []:

```
a = pred_df['Gpu'][0]  
pred_df['Gpu'] = gpudf.loc[gpudf.Gpu==a,'val'].values[0]
```

Linear Regression :

In []:

```
clf = LinearRegression ()  
clf.fit(X_train , y_train)
```

Out[]:

LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)

In []:

```
clf.score(X_train,y_train)
```

Out[]:

0.7963799723663094

In []:

```
clf.score(X_test,y_test)
```

Out[]:


0.7888777682574388


In []:

```
y_predicted = clf.predict(X_test)  
plt.scatter(y_test,y_predicted,color = '#88c999')
```

Out[]:

<matplotlib.collections.PathCollection at 0x7f19e57afe90>





Ridge:

In []:

```
ridlr = Ridge()
ridlr.fit(X_train, y_train)
score = ridlr.score(X_test, y_test)
score
```

Out[]:

0.7892888500291874

Random Forest Regressor :

In []:

```
from sklearn.ensemble import RandomForestRegressor
rfr = RandomForestRegressor()
```

In []:

```
rfr.fit(X_train, y_train)
```

Out[]:

```
RandomForestRegressor(bootstrap=True, ccp_alpha=0.0, criterion='mse',
                        max_depth=None, max_features='auto', max_leaf_nodes=None,
                        max_samples=None, min_impurity_decrease=0.0,
                        min_impurity_split=None, min_samples_leaf=1,
                        min_samples_split=2, min_weight_fraction_leaf=0.0,
                        n_estimators=100, n_jobs=None, oob_score=False,
                        random_state=None, verbose=0, warm_start=False)
```

In []:

```
rfr.score(X_train, y_train)  #train score
```

Out[]:

0.976429910826258

In []:

```
rfr.score(X_test, y_test)  #test score
```

Out[]:


0.8493155728100502


In []:

```
y_predicted = rfr.predict(X_test)
plt.scatter(y_test, y_predicted, color = '#88c999')
```

Out[]:

<matplotlib.collections.PathCollection at 0x7f19e5799ed0>





From the above, Random forest Regressor has the best Scores. So the required value to be predicted :

In []:

```
price = rfr.predict(pred_df)[0]  
price
```

Out[]:

308.47260000000001